

Session #03

Time Series Regression Models

What is a Time-Series Model?

- Regression applied to data indexed by time
- Captures: trends, seasonality, autocorrelation, external predictors
- Useful for understanding drivers & forecasting future values

Key Components

- Trend - Long-term direction
- Seasonality - Repeating patterns (monthly, yearly)
- Autocorrelation - Dependence on past values
- Exogenous Variables - Independent predictors (weather, ads, etc.)

Types of Time-series models

- Linear Regression with Time
- Autoregressive (AR) models
- Distributed Lag Models
- ARIMAX (ARIMA + Exogenous variables)
- Machine Learning Regression (RF, XGBoost, LSTM)

Building a Model (Steps)

- Explore Data - Plot, check ACF/PACF
- Feature Engineering - Add lags, seasonal dummies
- Model Training - Linear, ARIMAX, ML
- Diagnostics - Residual checks, error metrics
- Forecasting - Predict future values

Real-World examples

- **Retail:** daily sales using yesterday's sales, day-of-week, holidays, price, and promos.
- **Energy:** hourly demand using recent demand, hour of day, and temperature forecasts.
- **Healthcare:** ER visits using last week's counts, weekday patterns, and flu indicators.
- **Finance Ops:** call-center volume using month-end effects and product launch dates.

Same recipe, different domain.”

Key Takeaways

Time-series regression = regression that's time-aware.

It captures trend, seasonality, memory, and outside forces

